



DIRECTORATE OF RESEARCH, INNOVATION AND
PARTNERSHIP (DRIP), BAYERO UNIVERSITY, KANO

IMPERIAL

AI for literature review and academic writing

Shamsuddeen Hassan Muhammad

Bayero University, Kano, Imperial College London
Google DeepMind Fellow, Royal Society Career Fellow

<https://shmuhammadd.github.io>

BUK, Nigeria · 27 July 2026

1

IMPERIAL

About Me

- Google DeepMind Academic Fellow, Royal Society Career Development Fellow at Imperial College London
- Senior Lecturer, Bayero University, Kano · Visiting Lecturer, AIMS (Cameroon & South Africa)
- Co-founder, HausaNLP · Founder, Arewa Data Science Academy
- Research: Artificial Intelligence, Natural Language Processing, Low-resource African languages

2



IMPERIAL

2nd in the world
1st in the
UK and Europe

QS World
University
Rankings
2026


3

AI Research Began in the 1950s

- Alan Turing publishes his foundational paper in 1950
- The term “artificial intelligence” is coined at the 1956 Dartmouth Workshop
- Early machine learning programs emerge in the years that follow

The technology didn’t change overnight.


Three converging shifts explain why it exploded now.



4

Data, Compute and the 2017 Transformer Jointly Triggered the AI Boom


- **Data** — the web became a training corpus at trillion-token scale
- **Compute** — the cost of a unit of training fell by orders of magnitude
- **Algorithms** — the transformer architecture arrived in 2017
- **Take any one away** and none of this happens



Vaswani et al. (2017); Kaplan et al. (2020)

5

A few terms worth knowing



6


A few terms worth knowing

- 1 **Artificial Intelligence (AI)** — systems that perform tasks that normally need human intelligence, such as understanding language
- 2 **Machine Learning (ML)** — a subset of AI that improves at a task by learning patterns from data, rather than following fixed rules
- 3 **Large Language Model (LLM)** — a neural network trained on vast text data to predict and generate human-like language, e.g. GPT, Claude
- 4 **Generative AI** — AI that creates new content — text, images, code — rather than simply classifying or analysing existing data

7

Generative AI and Large Language Models (LLMs): What We Actually Use Today

- LLM is the text-focused subset of Generative AI:



8

Open and Close-Source GenAI Models

Two years ago there was one. Now there are dozens.

Closed Source

- ChatGPT
- Claude
- Gemini
- Copilot
- Grok
- Perplexity
- Amazon Nova

Open Source

- Llama
- Mistral
- DeepSeek
- Qwen
- Cohere Command
- Gemma
- Phi
- Kimi

Fifteen general-purpose models students can reach today, split by license — open-weight options now rival closed frontier models.

9

Open and Close-Source GenAI Models

Closed

- GPT, Claude, Gemini
- API access only; per-token cost
- Cannot be run on your own hardware
- Usually strongest at the frontier

Open

- Llama, Qwen, DeepSeek, Gemma
- Weights downloadable and inspectable
- Runs locally; no per-query cost
- Needs GPUs and expertise


10

GenAI in Research or Teaching

THE AI-ASSISTED RESEARCH PIPELINE



ADOPTION IS HIGH — TRANSPARENCY ISN'T



Sources: PLOS One, "Generative AI and academic scientists in US universities" (2025); Saito et al., Waseda University MBA thesis survey (Apr 2026); analysis of 5.2M papers, 2021–2025 (2026).

11

First AI System to Pass Peer Review at an A*-Ranked Venue

Tempest: Automatic Multi-Turn Jailbreaking of Large Language Models with Tree Search

Andy Zhou*
Intology AI
andy@intology.ai

Ron Arel*
Intology AI
ron@intology.ai

**ACL is the world's #1 scientific venue for Natural Language Processing (NLP), with an acceptance rate historically around 18–21%*

1 Introduction

Large language models (LLMs) have rapidly progressed in capability and accessibility, yet genuine safety validations often overlook how adversarial tactics can unfold across successive dialogue turns. While models are typically trained with safety constraints to refuse harmful requests, recent work has shown that these boundaries can be systematically eroded through repeated interactions (Li et al.,

¹We take responsibility for this work but the main intellectual contribution was conducted by an AI system

Source: Intology AI, Zochi paper acceptance announcement (ACL 2025); Zhou & Arel, "Tempest" (Intology AI, 2025).

12



Ravid Shwartz Ziv ✓ • 2nd

AI Researcher | Meta | NYU | Consultant | LLM...

1d •

[Connect](#)

NeurIPS reviews are out. At this point it's an arms race where both sides have the same supplier. You write the paper with Opus 4.6, the reviewer writes the review with Opus 4.7, and now we get to rebut with Opus 4.8. Let's just hope Opus 5 doesn't ship before the discussion period closes.



1,567



32



27




13

Fabricated Citations in Biomedical Papers Rose Twelvefold in Two Years

AI-hallucinated references are slipping past peer review at scale

- **4,406 fabricated references** found across 2,810 papers in a 97M-reference biomedical citation audit



Sources: Retraction Watch / Topaz et al. citation audit (2026); Ansari, "Compound Deception in Elite Peer Review", NeurIPS 2025 analysis (2026).

THE LANCET

This journal Journals Publish Clinical Global health Multimedia Events About

CORRESPONDENCE • Volume 407, Issue 10541, P1779-1781, May 09, 2026 [Download Full Issue](#)

Fabricated citations: an audit across 2.5 million biomedical papers

Maxim Topaz ^{a,b,c} • Nir Roguin ^{d,e} • Pallavi Gupta ^a • Zhihong Zhang ^{a,b} • Laura-Maria Peltonen ^{f,g,h}



[Affiliations & Notes](#) [Article Info](#) [Linked Articles \(3\)](#)

[Download PDF](#) [Cite](#) [Share](#) [Set Alert](#) [Get Rights](#) [Reprints](#)

14

Springer Nature retracted an entire book with fake citations

- Mastering Machine Learning: From Basics to Advanced — published April, retracted July 2025
- A fabricated citation was attributed to a real researcher, who noticed it



15

HKU Professor Steps Down After AI Citation Scandal

- HKU Professor Paul Yip stepped down as associate dean and withdrew from research committees after AI-fabricated citations were found in a paper he corresponding-authored

University of Hong Kong professor steps down from associate deanship after AI-generated references scandal

An internal probe by the University of Hong Kong confirmed that a paper – which listed Professor Paul Yip as corresponding author and was published in an academic journal – used non-existent publications generated by AI in its references.



Professor Paul Yip, director of the Hong Kong Jockey Club Centre for Suicide Research and Prevention at the University of Hong Kong, meets the press on September 10, 2025. Photo: Kelly Ho/HKFP.

Sources: Hong Kong Free Press (Dec 2025); VnExpress International (Dec 2025); Memon et al., Nature Human Behaviour (2025).

16

46% of Retracted Authors Leave Academic Publishing Entirely

- Retraction attaches to your name permanently and is indexed
- Institutions run misconduct proceedings; findings can end appointments

Local journals may not detect AI today

Sources: Memon et al., *Nature Human Behaviour* (2025); Retraction Watch (2025); Lu et al., *career effects of scandal study*.

17

Local Journals Often Lack the AI-Detection Tools International Publishers Use

- Bayero Journal of Pure and Applied Sciences (BAJOPAS)
- Bayero Journal of Engineering and Technology (BJET)
- Algaita: Journal of Current Research in Hausa Studies

Rarely have the budget for tools larger publishers use

18

AI is here to stay.

The responsibility now is learning to use it ethically.

19

Every Major Publisher permits AI use

All of them require you to declare it.

20

Check publishers page about the use of AI

Elsevier prohibit AI images:

The use of generative AI and AI-assisted tools in figures, images and artwork

Elsevier does not permit the use of generative AI or AI-assisted tools to create or alter images in submitted manuscripts. This may include enhancing, obscuring, moving, removing, or introducing a specific feature within an image or figure. Adjustments of brightness, contrast, or color balance are acceptable if they do not obscure or eliminate any information present in the original. Image forensics tools or specialized software might be applied to submitted manuscripts to identify suspected image irregularities.

<https://www.elsevier.com/en-gb/about/policies-and-standards/the-use-of-generative-ai-and-ai-assisted-technologies-in-writing-for-elsevier>

21

Every major publisher permits AI use and requires disclosure

Body	Position
ICMJE	No AI authorship. Writing assistance in acknowledgements; analysis in methods
COPE	AI cannot sign, cannot hold copyright, cannot take responsibility
Elsevier	Declaration in a separate section before the references; AI images prohibited
Springer Nature	Disclosure required; generative images prohibited
Wiley, T&F	AI cannot be an author; substantive use must be disclosed
SAGE	Assistive editing exempt; generative content must be disclosed and cited

Publisher author guidelines, July 2026. Check the specific journal — requirements vary within publishers.

22

Publishers Hold Authors, Not AI, Accountable for Papers

- A system cannot be held accountable — only a person can.
- “The AI wrote it” is not a defense.
- For supervisors, that includes your students’ submissions too.

COPE position statement (2023); ICMJE Recommendations, updated January 2026

23

Most publishers also forbid putting a manuscript you are reviewing into a chatbot

- A manuscript under review is confidential, and you are its custodian
- Pasting it into a commercial tool is a disclosure to a third party
- Several publishers prohibit this explicitly for editors and reviewers
- **The same applies to** grant applications, examiner reports and student theses

24

Don't outsource your thinking to AI.

Use it to search, structure and critique — never to think for you.

25

Good researchers get better with AI.

Researchers who know what they're doing get sharper.

Bad researchers get worse.

Those who just prompt and paste get dumber.

26

Skill With AI Tools Now Separates Colleagues' Outcomes

- 1 Colleagues who use these tools **well publish faster** and review better
- 2 Colleagues who avoid them **lose ground quietly**
- 3 Colleagues who use them **badly acquire retractions**

The differentiator is the skill. We need to learn to use AI tools to gain leverage.

27

Premium AI Model Costs Put Under-Resourced Researchers at a Disadvantage

- 1 Frontier plans charge steeply for full access — Claude Max runs \$100–\$200/month, ChatGPT Pro's \$200/month tier
- 2 Researchers and students without institutional funding must pay these costs individually, or rely on slower, rate-limited free tiers
- 3 This skews global competition: students at well-funded universities work with frontier models daily, while others fall further behind

Sources: Anthropic Claude pricing (2026); OpenAI ChatGPT Pro pricing (2026).

28

Some Universities Now Give Students Free Access to Frontier AI Models

- 1 Imperial's own dAlsy platform gives every student and staff member free access to ChatGPT, Claude, DeepSeek, Grok, and Llama
- 2 Oxford, LSE, Dartmouth, UC Berkeley, Duke, and the University of Chicago all now provide free institutional access to ChatGPT or Claude
- 3 Students without this kind of institutional deal are left choosing between rate-limited free tiers and subscriptions many can't afford

Where you study is starting to determine which tools you can use.

Sources: Felix Online, "Imperial grants all staff and students access to AI models" (2025); Chicago Maroon (2026); The Daily Texan (2026); Claude for Education partner list (2026).

29

May 30, 2024 Product

Introducing ChatGPT Edu

An affordable offering for universities to responsibly bring AI to campus.

Oxford gives students access to AI platform



IMPERIAL

- New Chat
- Artifacts
- Explore Agents
- Chats
- Agents

Previous 20 chats

Discover Intelligent Agents Designed by Our Team

Search for an Agent

Official Agents Engineered with permissions and safeguards for secure and controlled access

Browse all agents >

IMPERIAL Success Guide Chatbot

IMPERIAL's Success Guide chatbot, supporting students with learning, university life, and signposting to relevant Imperial...

Chat GPT-5.4

GPT-5.4 offers more accurate instruction-following, stronger multi-step reasoning, and clearer, more concise answers...

Claude Sonnet 4.5

Claude Sonnet 4.5 offers clearer, faster, more reliable responses with improved reasoning and instruction-following. Great for editing...

ChatGPT-5.3-codes with Code Interpreter

Focused on delivering performance for developers and professionals. Advanced coding, data analysis, and charting with code...

Chat GPT-5

Strengths: Cutting-edge reasoning, nuanced language understanding, strong multi-turn context retention, excels at complex analysis...

Chat GPT-5 (without live internet search)

Strengths: Cutting-edge reasoning, nuanced language understanding, strong multi-turn context retention, excels at complex analysis...

30

Universities Runs a Dedicated Hub Teaching Staff and Students to Use Generative AI

Generative AI & Education Guidance Hub

Welcome to the Generative AI and Education Hub, a specially developed resource for Imperial's education community to catch up on the latest developments in the use of generative artificial intelligence across our learning environments.


For staff and students, these generative AI models present both opportunities to further improve the quality of our teaching and assessment, as well as various challenges. Knowledge, transparency, and innovation will be key principles driving Imperial's successful interaction with generative AI, as well as debates on and off our campuses about how best it can be applied to real-world scenarios.

Are you a student? Visit the Generative AI and Study Guidance Hub >

Have you explored Imperial's Vision for Education in the Age of AI? >



Introduction to Generative AI



Generative AI Principles



Networks and Events

As generative AI guidance is further developed and more

Universities worldwide — including Oxford, LSE, Dartmouth and UC Berkeley — now run similar staff and student AI-literacy programmes.

31

AI vs Researcher

- AI makes a careful researcher faster and an uncritical one dangerous
- Learn the craft first — the tools amplify whatever is already there
- The people who do best are those who can tell when the model is wrong

32

Keep these five commitments in mind whenever you use AI

- Use AI to search, structure, critique and edit — never to assert
- Never cite a paper you haven't opened
- Never publish a sentence you couldn't defend under questioning
- Never paste confidential work into a tool you don't control
- Disclose what you used, and where

33

Prompting: The One Skill That Makes AI Actually Useful

The next few slides cover the one skill that determines whether AI helps you or wastes your time.

34

Prompting Is the Core Skill for Getting Excellent AI Output

The same model gives dramatically different answers depending on how you ask.

- A vague prompt gets a generic, often unusable answer.
- A well-crafted prompt gets a precise, expert-level answer — from the exact same model.
- The difference is entirely in how you ask, not which tool you use.
- This is the single highest-leverage skill in using AI well. Everything else — literature review, data analysis, writing — builds on it.

Next: the four elements that make a prompt excellent.

35

A Good Prompt Has Four Elements

Role, Context, Task, Format — miss one and the model has to guess

- **Role** — tell it who to be — “You are a senior structural engineer.”
- **Context** — give it the background it actually needs — the facts, constraints, and audience
- **Task** — state exactly what you want done, not just the topic
- **Format** — specify the structure, length, and style of the output you want back

Rule of thumb: The more specific your prompt, the more useful the output. Vague prompts yield vague results.

36

Example: A Prompt Built for a Doctor

Differential diagnosis: same question, two very different outputs

Vague prompt: "What are symptoms of diabetes?"

Excellent prompt:

"You are an endocrinologist. Summarize the differential diagnosis for a 45-year-old patient presenting with polyuria, unexplained weight loss, and fatigue over 3 months. Include the key lab tests to order first, ranked by urgency. Format as a clinical triage table."

Role: "You are an endocrinologist."

Context: "A 45-year-old patient presenting with polyuria, unexplained weight loss, and fatigue over 3 months."

Task: "Summarize the differential diagnosis and include the key lab tests to order first, ranked by urgency."

Format: "Present as a clinical triage table."

37

A vague prompt returns generic textbook symptoms with no patient-specific triage guidance.

Vague prompt used: "What are symptoms of diabetes?"

Output:

No role, no patient details, no requested format — the model defaults to a generic textbook answer with no clinical prioritization.

What are symptoms of diabetes?

Here are the common symptoms of diabetes:

Classic symptoms (both type 1 and type 2):

- Frequent urination (polyuria)
- Excessive thirst (polydipsia)
- Increased hunger, even while eating
- Unexplained weight loss
- Fatigue and tiredness
- Blurred vision
- Slow-healing sores or cuts
- Frequent infections (skin, gums, bladder, vaginal)
- Tingling or numbness in hands or feet

Type 1 diabetes specific notes:

- Symptoms tend to develop quickly, over days or weeks
- Can include nausea, vomiting, and stomach pain
- Often diagnosed in children, teens, or young adults, though it can occur at any age

Type 2 diabetes specific notes:

- Symptoms often develop slowly, sometimes over years, and can be mild enough to go unnoticed at first
- Some people have no symptoms at all and are diagnosed through routine blood work
- Darkened skin patches (acanthosis nigricans), often around the neck or armpits, can be a sign of insulin resistance

Signs of a diabetes-related emergency (seek immediate medical attention):

- Diabetic ketoacidosis (mainly type 1): fruity-smelling breath, rapid breathing, confusion, vomiting
- Severe hypoglycemia (low blood sugar): shakiness, sweating, confusion, seizures, loss of consciousness
- Hyperosmolar hyperglycemic state (mainly type 2): extreme thirst, very high blood sugar, confusion, in some cases coma

38

A structured prompt returns a ranked, patient-specific differential ready for clinical triage.

Excellent prompt used: "You are an endocrinologist... 45-year-old patient, polyuria, weight loss, fatigue... triage table."

Output:

Same underlying question, now scoped by role, patient context, and output format — the model returns a ranked, actionable differential.

Rank	Suspected Diagnosis	Supporting Features	Key First-Line Labs	Urgency
1	Diabetes mellitus (T1/T2) ± DKA/HHS	Classic triad; osmotic diuresis drives polyuria/weight loss	Random/fasting glucose, HbA1c, venous blood gas, serum + urine ketones, basic metabolic panel	Immediate — rule out acute hyperglycemic crisis
2	Hyperthyroidism	Weight loss despite appetite; fatigue, heat intolerance, tremor	TSH, free T4, free T3; TRAb if positive	High
3	Diabetes insipidus (central/nephrogenic)	Marked polyuria/polydipsia with normal glucose	Serum + urine osmolality, serum Na ⁺ , serum glucose (to exclude DM first)	High
4	Hypercalcemia (e.g., hyperparathyroidism, malignancy)	Polyuria, fatigue, weight loss, "stones, bones, groans"	Serum calcium (corrected/ionized), PTH, phosphate	High
5	Adrenal insufficiency	Fatigue, weight loss; hyponatremia, hyperkalemia	Morning cortisol, ACTH, electrolytes	Moderate
6	Occult malignancy	Unexplained weight loss, fatigue; paraneoplastic effects	CBC, ESR, CRP, LFTs, age-appropriate imaging/screening	Moderate — pursue if endocrine workup negative
7	Chronic infection (e.g., TB, HIV)	Constitutional symptoms, weight loss	CBC, HIV serology, targeted screening	Moderate

First-tier bundle to draw simultaneously: glucose + HbA1c, TSH/free T4, corrected calcium, electrolytes/renal panel, CBC. This single panel screens the top diagnoses and flags any needing urgent intervention.

Red flags for immediate escalation: altered mental status, dehydration, glucose >250 with ketones, or Kussmaul breathing → treat as DKA/HHS pending confirmation.

39

Example: A Prompt Built for an Economist

Policy briefing: same question, two very different outputs

Vague prompt: "Explain inflation."

Excellent prompt: "You are a macroeconomist advising a central bank. Explain how a 2-point interest rate hike would likely affect inflation, unemployment, and currency value in an import-dependent economy over 12 months. Cite the key transmission mechanisms and flag major uncertainties. Format as a one-page briefing memo for a non-economist minister."

Role: "You are a macroeconomist advising a central bank."

Context: "A 2-point interest rate hike in an import-dependent economy, assessed over a 12-month horizon."

Task: "Explain the likely effects on inflation, unemployment, and currency value. Cite the key transmission mechanisms and flag major uncertainties."

Format: "Present as a one-page briefing memo for a non-economist minister."

40

Literature Review with AI

41

AI can help with 4 of the six literature-review stages, but cannot automate 2.

Across the six-stage workflow, four stages can run with AI assistance; two still require a human in the

1. Formulate a searchable question — AI-assisted
2. Search & discover the literature — AI-assisted
3. Screen & select relevant studies — Human required
4. Extract & organize study data — AI-assisted
5. Appraise evidence quality & verify sources — Human required
6. Draft & structure the narrative — AI-assisted

42

AI Tools for Literature Review

A new ecosystem of AI-powered tools is transforming how researchers discover, read, and synthesize literature. Here's what you need to know.

Discovery & Search

- **Consensus** — AI-powered academic search
- **Elicit** — research assistant for paper discovery
- **Semantic Scholar** — AI-enhanced literature search

Synthesis & Writing

- **ChatGPT / Claude** — general synthesis and drafting
- **Perplexity** — AI search with citations
- **ResearchRabbit** — visual literature mapping

Reading & Summarization

- **ChatPDF / Claude** — chat with PDFs
- **NotebookLM** — Google AI for document synthesis
- **SciSpace** — AI copilot for reading papers

Local & Private Options

- **LM Studio** — run open models locally
- **Ollama** — local LLM runner
- **AnythingLLM** — private RAG for your papers

Golden rule: No single tool does everything. Build a toolchain — use each tool for what it does best.

43

The Research Tool Landscape

Tool	Type	Strength	Best For
ChatGPT	General LLM	Versatile, fast, good at synthesis	Drafting, brainstorming, explaining
Claude	General LLM	Long context, nuanced reasoning	Deep document analysis, long PDFs
Perplexity	AI Search	Real-time web search + citations	Fact-checking, current events
Consensus	Academic Search	Academic-focused, evidence summaries	Finding papers, consensus mapping
NotebookLM	RAG Platform	Private, sources-grounded answers	Synthesizing your own document set
ResearchRabbit	Visual Discovery	Citation networks, visual maps	Exploring related literature
Elicit	Research Assistant	Structured extraction from papers	Systematic reviews, data extraction

Landscape as of mid-2026.

44

No single tool covers the workflow

Tool	Best for	Cost
Semantic Scholar	Discovery across 200m papers; citation graph	Free
Connected Papers	Mapping a subfield from a seed paper	Free, 5 graphs/month
Elicit	Screening and structured extraction at scale	Free tier; ~\$12/month
Consensus	Fast directional evidence check	Free tier
NotebookLM	Grounded synthesis over your own PDFs	Free tier
Zotero	Citation management and bibliography	Free
Claude / ChatGPT	Reasoning, critique, revision — never citations	Free tier

45



A scholarly research assistant that combines literature understanding and data-driven discovery. Asta uses 108M+ abstracts and 12M+ full-text papers to find, summarize, and analyze scientific evidence. A project from [AI2](#).

[Find papers](#)
[Generate a report](#)
[Analyze data](#)
 Describe the papers you are looking for

Deep search ▾ ➤

[Show example queries](#)
[Explore research prototypes in ✨AstaLabs](#)

46

Step 1: Discovery — Finding the Right Papers

AI can help you discover papers you didn't know to look for.

Method A — Semantic Search

Describe your research question in natural language. AI tools find conceptually related papers even when keywords don't match.

Example: "papers about using machine learning to predict crop yields in arid regions"

Method B — Citation Chaining

Start with one key paper. Use ResearchRabbit or Semantic Scholar to explore who cited it, what it cited, and related work.

Visual citation maps reveal research clusters you might miss otherwise.

SAMPLE PROMPT

"I am researching [TOPIC]. Find me the 10 most influential papers from the last 5 years. For each, explain: (1) the core contribution, (2) the methodology, (3) how it relates to my topic, and (4) one limitation I should be aware of."

Always verify: AI can hallucinate paper titles. Cross-check on Google Scholar before citing.

47

Step 2: Synthesis — Making Sense of What You Found

. The real skill is synthesizing papers into a coherent narrative that identifies gaps, conflicts, and opportunities.

What AI Can Do

- Summarize individual papers in plain language
- Compare and contrast methodologies across studies
- Identify recurring themes and conflicting findings
- Generate tables of study characteristics
- Suggest gap areas based on what's missing
- Draft section outlines for your review

What ONLY You Can Do

- Judge the **quality** and credibility of sources
- Recognize **subtle methodological flaws** AI misses
- Identify **your unique angle** — what hasn't been asked
- Understand **disciplinary context** and politics
- Make **normative judgments** about what matters
- Craft a **compelling narrative** for your audience

SAMPLE PROMPT

"I have found these 5 papers on [TOPIC]. Synthesize them into a narrative that shows: how the field evolved, where researchers agree, where they disagree, and what gap my work could fill."

48

Step 3: Critical Analysis — Don't Trust, Verify

AI can summarize, but it cannot think. Critical analysis is where your expertise becomes irreplaceable. Here's a checklist.

1. Verify citations

AI invents references. Every citation AI suggests must be checked on Google Scholar or PubMed before inclusion.

2. Check methodology

Does the study design actually support the conclusion? AI often misses weak causal claims.

3. Assess sample and scope

Is the finding generalizable? AI summaries often strip away the "in this specific population" qualifiers.

4. Look for conflicts of interest

Funding sources and author affiliations matter. AI rarely flags these unless explicitly asked.

Remember: AI is your research assistant, not your replacement. Your critical judgment is the final filter.

49

Academic writing

Structure, critique and edit. Do not generate.

50

Use AI where it is strong — structuring, critiquing and editing

Works well

- Outlining and reverse-outlining
- Reviewer simulation before submission
- Revision for concision and clarity
- Language editing and register

Works badly

- Generating claims you have not evidenced
- Producing reference lists
- Writing your contribution for you
- Assessing novelty against the literature

51

Five Prompt anatomy with Example

Five components. Most weak prompts are missing three of them.

ROLE	You are a reviewer for [venue], expert in [subfield].
CONTEXT	This is the [section] of a paper on [topic]. The audience is [who].
TASK	[One instruction, stated as an imperative.]
CONSTRAINTS	Keep it under [N] words. Preserve my technical claims exactly. Do not add citations. Flag any claim I have not supported.
EXAMPLE	Here is a paragraph of mine you should match in tone: [paste].

52

Asking the model to attack your draft as Reviewer 2

You are a sceptical reviewer for [venue]. You are not hostile, but you are unconvinced, and you have read this literature closely.

Write the three strongest objections to this paper. For each: state the objection, say precisely what evidence would answer it, and say whether the paper as written contains that evidence.

Do not comment on writing quality. Do not be encouraging.

Then answer the objections in the paper, before a real reviewer raises them.

53

Disclosure

STANDARD FORM

"During the preparation of this manuscript the author(s) used [tool, version, publisher] in order to [specific purpose]. The author(s) reviewed and edited the output as necessary and take full responsibility for the content of the publication."

WHERE IT GOES

Writing assistance → acknowledgements. Data collection or analysis → methods.
Elsevier: a separate declared section immediately before the references.

Keep a log of what you used and where. If an editor queries it, you will want the record.

54

Conclusions

- 1 **Generative models produce text, not truth.** Every citation must be verified against a DOI before it enters a draft.
- 2 **Fabricated references have risen twelve-fold in three years.** Screening is arriving, and the record is permanent.
- 3 **Responsibility never transfers.** You sign the paper; you own every sentence in it, including your students'.
- 4 **Use AI to search, structure, critique and edit.** Never to assert. Arrangement can be assisted; assertion cannot.

55

References

- ACL Rolling Review (2026). Cycle statistics, January and May 2026 cycles. aclrollingreview.org
- Ansari (2026). Compound deception in elite peer review: a failure mode taxonomy of 100 fabricated citations at NeurIPS 2025. [arXiv:2602.05930](https://arxiv.org/abs/2602.05930).
- Committee on Publication Ethics (2023). Authorship and AI tools: COPE position statement.
- Elsevier (2025). The use of generative AI and AI-assisted technologies in writing for Elsevier.
- ICMJE (2026). Recommendations for the conduct, reporting, editing and publication of scholarly work in medical journals.
- Imperial College London (2026). Generative AI guidance; Code of Practice 8; dAIsy platform documentation.
- Kaplan et al. (2020). Scaling laws for neural language models. [arXiv:2001.08361](https://arxiv.org/abs/2001.08361).
- Liang, Yuksekgonul, Mao, Wu & Zou (2023). GPT detectors are biased against non-native English writers. *Patterns*, 4(7), 100779.
- Liang et al. (2024). Monitoring AI-modified content at scale: the impact of ChatGPT on AI conference peer reviews. *ICML, PMLR* 235.
- NeurIPS Program Committee Chairs (2025). Reflections on the 2025 review process. blog.neurips.cc
- Topaz et al. (2026). Fabricated references in PubMed-indexed articles. *Correspondence, The Lancet*.
- Vaswani et al. (2017). Attention is all you need. *NeurIPS*.
- Reporting on the distillation allegations: CNBC and CNN, 24 February 2026.

56

Thank you

shmuhammad.csc@buk.edu.ng
s.muhammad@imperial.ac.uk

<https://shmuhammadd.github.io>